

Layer 1 is basically okay. I would only make **two small corrections** before moving on.

Small Layer 1 corrections

1. prepare_runpod_workspace should know about project namespaces

Layer 1 should not create research code, but it should prepare storage roots that later bundles can use consistently:

```
/workspace/repos/  
/workspace/data/  
/workspace/runs/  
/workspace/artifacts/  
/workspace/models/  
/workspace/checkpoints/
```

Then later bundles create project-specific paths under those roots:

```
/workspace/repos/nca-art-grn  
/workspace/data/nca-art-grn  
/workspace/runs/nca-art-grn  
/workspace/artifacts/nca-art-grn
```

So Layer 1 remains generic. Layer 2/3 create the research-specific subfolders.

2. Remote model endpoint belongs in Layer 1, but only as a thin contract

The remote model loop should stay minimal:

```
local code -> remote model -> response
```

No Agentfield, no Paperclip, no OpenClaw logic yet. The uploaded remote-model plan explicitly says the starting point is one reliable loop, where Runpod serves the model and local code talks to it through a standard chat API.

So Layer 1 is good if we keep it as:

machine/runtime paths
GPU/CUDA checks
Docker/runtime policy
remote model client contract

Nothing else needs changing